

ASHRITHA GONUGUNTLA

agonugun@andrew.cmu.edu | +1 (412) 209-9824 | ashrithag.github.io | linkedin.com/in/ashrithagonuguntla | github.com/AshrithaG | scholar

Pittsburgh, PA · Available January 2027

EDUCATION

Carnegie Mellon University, School of Computer Science

Pittsburgh, PA

Master of Software Engineering, **GPA 4.0/4.0**

Aug 2025 – Dec 2026

- Coursework: AI Agents, Deep Learning Systems, Generative AI, Multimodal ML, Diffusion and Flow Matching, ML in Production, Computer Systems, Software Architecture, Mobile Robots, F1-Tenth Autonomous Racing.
- Teaching Assistant, four courses: ML in Production (Prof. Kästner), Prompt Engineering (Prof. Breaux), Computer Systems (Prof. Kesden), Engineering Computation (Prof. Yamakawa).

K L University

Andhra Pradesh, India

B.Tech Computer Science, AI & Robotic Process Automation, **GPA 9.84/10, rank 2**

Jun 2018 – Mar 2022

- Academic Excellence Award for second rank in the graduating class.

EXPERIENCE

eParts Services (CMU MSE Studio client), Technical Lead, ML & Architecture

2026 – Present

- Leading **5 engineers** shipping a confidence-routed attribute-prediction service to the client: BGE+FAISS retrieval, a tiered rule engine, and Mahalanobis confidence routing over 4.08M attribute rows; 95.74% ProductType accuracy on 12,958 held-out products.
- Caught a deployment-blocking evaluation error: semantic-only scoring capped fused confidence at 0.30, making the **0.85 auto-accept gate mathematically unreachable**; added end-to-end rule composition and per-ProductType calibration before rollout.
- Hardened the Python 3.12/FastAPI service with ruff, black, mypy, **289 tests, and an 85% branch-coverage gate**; load-tested Prometheus-instrumented serving at 50 req/s against a 200 ms p95 SLO and isolated p95 latency to the encoder (~38 ms).

Cisco, Software Engineer → Senior Software Engineer

Aug 2022 – Jul 2025

- Tech-led 9 engineers across Cisco Secure Firewall's two telemetry paths in C++/Go: a gRPC plane with control/data channels, protobuf metrics, mTLS, ACK-based cancellation, plus gNMI/OpenConfig YANG subscriptions, scale-tested to 1,000 collector endpoints; the feature carries **\$100M in ARR**.
- Rebuilt the intrusion-event stream-processing and query path with tailored indexing and range-based partitioning: **40% lower latency at 2x telemetry volume**; mentored 4 engineers through design and performance reviews.
- Cut median investigation-preparation time from **30 to 21 minutes** in a 4-week pilot with an agentic RAG incident assistant built with 8 SOC analysts: multi-step LLM orchestration and tool calling over firewall telemetry and configuration, **Amazon Bedrock** inference, SageMaker embeddings and OpenSearch retrieval on AWS, with diagnostic tools kept read-only.
- Contributed to **LLM post-training** for a 1B security model, later released as **Antares-1B**: prepared CVE/CWE SFT data, evaluated early checkpoints, built network-disabled Docker sandboxes for read-only repo exploration; scope preceded GRPO.
- Carried **production on-call** for telemetry and observability services: alert triage and log/query diagnosis through mitigation and follow-up fixes across ingestion and operator-facing detection workflows.

Cisco, Software Intern

Jan 2022 – Jun 2022

- Built cloud-native security observability on Kubernetes with **Prometheus, Grafana, CloudWatch**; automated AWS provisioning with Terraform, Ansible, CloudFormation.

Adobe Research, Research Intern

May 2021 – Aug 2021

- Co-authored **Outage-Watch** (ACM ESEC/FSE'23), a mixture density network for QoS tail-risk forecasting: **AUC 0.98, MTTD cut up to 88%**; co-inventor on **US Patent App. 17/656,263**.

IIIT Hyderabad, Research Intern, Speech and Vision Lab

Apr 2021 – May 2021

- Implemented neural architectures for automatic speech recognition and text-to-speech synthesis in **TensorFlow and PyTorch**.

University of Kerala, ML Intern, Computational Biology & Bioinformatics

Oct 2020 – Dec 2020

- Built a **CycleGAN** pipeline for cross-domain style transfer, enhancing structural patterns in cellular microscopy images.

Microsoft, Software Engineering Intern

May 2020 – Aug 2020

- Built the metrics and insights layer of **Viva Insights** over terabytes of data for 27K+ companies, in Azure and **PySpark**; validated the influence and network algorithms end to end for customer release.

PUBLICATIONS

- *The Replay Gap: Static Evaluation of Model Switching in LLM Agents Scores the Wrong World* (sole author). **COLM'26 Workshop on Efficient Reasoning, Spotlight**. Branching rollouts over ~900 SWE-bench trajectories show static replay misscores 92–97% of switching decisions; serving-stack cost and determinism characterized over 20,141 requests.
- *The Plan, Not the Decoder: Diagnosing and Repairing Compositional Failure in Reasoning-Augmented Text-to-Image Generation* (sole author). **ECCV'26 MUCG Workshop, Oral**. Localized compositional failure to the planner rather than the decoder; training-free repair gained +10.7 points over ~12K images.
- *APIShift: Verifiable Agent Learning for API Contract Migration*. **EMNLP 2026, under review**. A verifiable RL environment over 1,549 real API version pairs from 449 providers; GRPO and LoRA training on Qwen2.5-3B with a documented reward-variance collapse and its repair directions.

PROJECTS

int8-linear | *Triton W8A8, vLLM 0.28, CUTLASS baseline, Qwen3-1.7B, RTX 4090*

[Code](#)

- Wrote a fused **W8A8 int8 Triton kernel** and served Qwen3-1.7B end to end in **vLLM** under CUDA graphs: decode reaches **1.15x to 1.19x the tokens per second** of vLLM's default CUTLASS kernel at batch 32–128 and 1.29x on prefill, at **no measurable accuracy cost** (WikiText-2 perplexity 20.52 against 20.49, identical logprobs at all 20,440 tokens across three runs).
- Predicted from vLLM's sm89 dispatch source that CUTLASS would step at a bucket edge, then swept and found it: **1.31x to 1.99x slower from M=16 to M=17** on five of six layer shapes, and 1.06x on the sixth, which takes a different configuration there. Reported upstream as [vllm-project/vllm#56924](#).

nanoinfer | *C++20, ARM NEON, CUDA, Triton, cuBLAS/cuDNN, ONNX Runtime*

[Code](#)

- Built a C++20 edge inference engine (im2col+SGEMM, NEON depthwise, fused epilogues, liveness-planned arena with no forward-pass allocations): **0.99x ONNX Runtime** on the dense CNN, 0.60–0.77x on depthwise, 37–77% less activation memory.

replay-gap | *vLLM, FP8/AWQ, SWE-bench, mini-SWE-agent*

COLM'26 Workshop Spotlight, sole author

- Built a resumable branching-rollout harness that forks live SWE-bench agent trajectories, rebuilds each Docker environment, and continues the fork under a different model: **889 rollouts** over 20,141 requests and 42.5 GPU-hours.

APIShift | *GRPO, LoRA, Qwen2.5-3B, verifiable rewards, RTX 3090*

EMNLP 2026, under review

- Built a verifiable RL environment for API contract migration over **1,549 real API version pairs** from 449 providers and ten breaking-change types: a manager emits JSON actions to Diff, Patch, Test and Rollback specialists under a five-component reward.

batch-invariance | *vLLM 0.28, Qwen3-1.7B, RTX 4090, CUDA graphs*

[Code](#)

- Built a probe harness isolating batch size, batch composition, repeat rate, and throughput against vLLM 0.28, and priced the determinism fix at **54 to 67%** of throughput, a cost the documentation states but never quantifies.

specdec | *PyTorch, Qwen3-0.6B draft, Qwen3-4B target, RTX 4090*

[Code](#)

- Implemented speculative decoding from scratch and showed it **loses at every window size** in eager PyTorch (0.87x to 0.51x): the draft costs 0.649 of a target pass, so break-even reduces to acceptance = cost ratio.

Query-Rewrite Evaluation | *Qwen3-1.7B, LoRA (PEFT), BM25, FastAPI, QReCC, nDCG, Kendall τ*

[Code](#)

- Fine-tuned Qwen3-1.7B with LoRA (PEFT, prompt-masked loss, cosine schedule) on 48,415 QReCC pairs, lifting nDCG@10 from 0.400 to **0.481** against a 0.505 human reference; saved 7 checkpoints as evaluation systems.

Gated Residual RL | *MuJoCo, MetaWorld, PPO, SAC, PETS/MPPI*

[Code](#)

- Caught a policy scoring **85% in metrics but 0% on camera**, traced it to seed and protocol leakage, and corrected the evaluation before reporting.

mlops-replay | *Python, paired bootstrap, drift monitoring, Docker*

[Code](#)

- Replayed **120 months** of production-style data through score/label/retrain/gate/promote; a paired-bootstrap gate shipped 3 regressions vs. 22 for monthly retraining.

F1TENTH Autonomous Racing | *ROS2 Humble, Python, Jetson, Hokuyo LiDAR*

CMU 16-663, team

- Built reactive navigation on a physical 1/10-scale racecar: iTTC emergency braking, PID wall-follow, disparity-extended Follow-the-Gap, adaptive Pure Pursuit, RRT*. **Won 3 of 3 races.**

SiMATU: LLM Unlearning | *Phi-3.5-mini 3.8B, RWKV, causal tracing*

CMU Deep Learning Systems, team

- Made a 3.8B model forget targeted facts: knowledge sits in early layers but only late-layer edits remove it, so a three-layer edit hit **91.9% unlearning** at 97.4% retention, up from 74.1%.

Autonomous Robotic Arm for EV Charging | *computer vision, 6-DOF manipulation, team lead*

K L University

- Led a team building the perception pipeline for a **6-DOF robotic arm** performing autonomous EV charger docking: object detection and spatial reasoning to localise the charging port.

HONORS AND AWARDS

- Mary Shaw Award for Creativity in Software Engineering**, Carnegie Mellon, 2026 · 5th place, Gray Swan AI red-teaming challenge, 2026 · CMU Moon Miners, NASA Lunabotics, 2026.
- JN Tata Scholar** and **KC Mahindra Scholar**, 2025 · 2nd place, Best Adoption category, AWS Graviton Challenge, 2024 · Recognised under Cisco's Achieve Operational Excellence category.
- Winner, Adobe SheCodes**, 2020 · **National winner, Smart India Hackathon**, 2020 · McKinsey Next Generation Women Leaders, 2021 · Walmart CodeHers finalist, 9th rank, 2021 · Cisco Ideathon finalist, 2021.
- Gold Medalist and Best Outgoing Student**, CBSE secondary school board, class of 2015.

LEADERSHIP AND SERVICE

- Regional Chapter Lead, GirlScript Foundation India**: led the Andhra Pradesh chapter of a nonprofit running free coding programs for women and underrepresented students.
- Volunteer, eVidyaloka Trust and Empower Ananya**: digital literacy for students in underserved rural and urban communities across India, through video classrooms and outreach.
- Honored Speaker, ShaktiCon 2023**: “5G IoT Security: Protecting the Next Generation of Connected Devices.” · MSE Leadership Initiative, Carnegie Mellon.
- Co-founder, Vector AI** (2019–2020): a cultural gaming startup; shipped *Ethan's Adventure*, a puzzle game with custom enemy and character AI, to the Google Play Store.

TECHNICAL SKILLS

Languages: Python, C++20, Go, Java, C, SQL, Bash, TypeScript/JavaScript.

Additional tooling: TensorFlow, scikit-learn, Ray, Flower, MLflow, PySpark, Redis, MongoDB, Oracle, Splunk, Helm, Terraform, Ansible, GCP, Azure, Isaac Sim, Gazebo, Nav2, GoogleTest.